Vaishnav Potlapalli

332-373-5002 vp2499@nyu.edu linkedin.com/in/vaishnav-potlapalli github.com/va1shn9v

Publications

PromptIR: Prompting for All-in-One Blind Image Restoration **NeurIPS 2023**

Vaishnav Potlapalli*, Syed Waqas Zamir, Salman Khan, Fahad Shahbaz Khan*

- Proposed an implicit prompt-learning based approach for All-in-One blind Image Restoration. Achieves SoTA performance on multiple image restoration tasks,without any prior degradation information.

Sketch3T: Test-Time Training for Zero-Shot SBIR **CVPR 2022**

Aneeshan Sain, Ayan Kumar Bhunia, ***Vaishnav Potlapalli****, Pinaki Nath Chowdhury, Tao Xiang, Yi-Zhe Song*

- Introduced a novel test-time training paradigm for zero-shot sketch-based image retrieval that adapts to new categories and sketch distributions using a single sketch, outperforming state-of-the-art methods.

MediTables IIIT **GREC 2021**

Akshay Praveen Deshpande, ***Vaishnav Potlapalli****,Ravi Kiran Sarvadevabhatla*

- Built a new dataset and semantic segmentation model for camera captured medical document images.

Experience

Floma Inc - AI Agents for Marketing **January 2025 – Present**

Founding Software Engineer - Machine Learning

- Engineered a CV auto-labeling pipeline fusing OW-Detection (Grounding-DINO) with SAM segmentation, achieving a **5×** dataset-annotation speedup.
- Fine-tuned Llama-family models on ad-copy data via **SFT /RLHF**, boosting generation performance by **23%**.
- Built a **multimodal LLM agent** that programmatically composes dynamic display ads; created an SVG generation/editing module for size-agnostic ad rendering and on-the-fly customization.

MBZ University of Artificial Intelligence **July 2022 – July 2023**

Research Assitant - Computer Vision *(Advisor: Dr. Salman Khan)*

- Proposed and implemented a novel Visual transformer based prompt-learning framework for All-in-one blind Image Restoration / Enhancement called **PromptIR**,which achieved **SoTA** performance improving over previous methods by **0.9 dB** on dehazing, deraining and denoising benchmarks.Work presented as part of **Neurips 2023**
- Adapted computer vision based continual learning techniques **L2P**, **DualPrompt** methods for video action recognition improving performance over previous techniques by over **10%** accuracy and **14% BWF**, on several public benchmarks.
- Studied parameter-efficient finetuning techniques to improve downstream performance of **Multimodel LLM models**.

Dhan AI **April 2020 – April 2022**

Machine learning Engineer

- Developed an ensemble of **BERT-based Classifiers** to enhance the NER engine, resulting in a **12% accuracy** improvement in internal benchmarks on entity recognition and sentiment classification, significantly improving the primary product of the company, which was a Patient Life Cycle Managment Chatbot.
- Rewrote the application testing pipeline to utilize increased **parallelism and Redis cache** to reduce CI/CD time by **60%** that enabled rapid development of new features.

Education

New York University Courant Institute of Mathematical Sciences **Sep. 2023 – Dec 2024**

Masters of Science in Computing, Entrepreneurship, and Innovation GPA: 3.96/4.0 *New York City, NY*

Relevant Courses*: LLVMs, Big Data and ML Systems, Foundations of Computer Networks*

Honors*: M. Michael Waller Master's Fellowship*

Projects

Efficient Mixture-of-Depths(MoD) LLM Inference | *PyTorch, CUDA, LLMs* **March 2024 – May 2024**

- Engineered a Mixture-of-Depths (MoD) transformer on a LLaMa-style baseline (55M parameters, 6 layers) by integrating dynamic token routing with top-k selection, auxiliary loss, and an auxiliary MLP predictor
- Utilized Torch CUDA events for profiling and simulated a novel GPU scheduling policy to boost throughput, lower latency, and validate improvements via ablation studies and perplexity analysis.

PITCHPAL: AI-Powered Presentation Coach | *FastAPI, React, LLMs, AI Agents* **October 2024 – December 2024**

- Implemented a multi-modal AI Agent with a fast-api backend that combined advanced speech recognition, natural language processing, and computer vision to analyze presentation content and delivery and provide relevant feedback.

Technical Skills

Languages/Frameworks: Python, Java, C/C++, CUDA, TypeScript, Pytorch, TF,OpenCV, Transformers, Django, MySQL